

DATA WRANGLING PROJECT THREE

```
#  
=====  
=====  
# Step 1: Setup and Data Loading  
#  
=====  
=====  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from scipy import stats  
from sklearn.model_selection import train_test_split  
  
# Setting the visual style for plots  
sns.set_style("whitegrid")  
  
# Loading all the datasets from CSV files  
try:  
    invoices_df = pd.read_csv("synthetic_invoices.csv")  
    vials_df = pd.read_csv("synthetic_vials.csv")  
    dispense_log_df = pd.read_csv("synthetic_dispense_log.csv")  
    claims_df = pd.read_csv("synthetic_claims.csv")  
    print("All CSV files loaded successfully.")  
except FileNotFoundError as e:  
    print(f"Error loading data: {e}. Please ensure all CSV files are  
in the correct directory.")  
    raise  
  
All CSV files loaded successfully.  
  
#  
=====  
=====  
# Step 2: Data Wrangling (Cleaning and Merging)  
#  
=====  
=====  
# The 'Purchase Price' column in vials_df to be a numeric type because  
# they have '$' in it.  
vials_df['Purchase Price'] = vials_df['Purchase Price'].replace({'\\$': ''}, regex=True).astype(float)  
  
# The separated other csv files merging into a single dataframe,  
# comprehensive dataframe  
vial_invoice_df = pd.merge(vials_df, invoices_df, on="Invoice Number",  
how="left")
```

```

full_df = pd.merge(vial_invoice_df, dispense_log_df, on="Vial Number",
how="left")
full_df = pd.merge(full_df, claims_df, on="Dispense ID", how="left")
print("\nDataframes merged into 'full_df'.")
```

Dataframes merged into 'full_df'.

```

# Checking for and to handle duplicate rows
duplicates = full_df.duplicated().sum()
print(f"\nNumber of duplicate rows found: {duplicates}")
if duplicates > 0:
    full_df = full_df.drop_duplicates()
    print("Duplicates have been removed.")
else:
    print("No duplicates found.")
```

Number of duplicate rows found: 0
No duplicates found.

```

# Checking for and to handle missing values using the forward-fill method
print(f"\nTotal missing values before handling:
{full_df.isnull().sum().sum()}")
full_df.fillna(method='ffill', inplace=True)
print("Missing values handled using forward fill.")
print(f"Total missing values after handling:
{full_df.isnull().sum().sum()}")
```

Total missing values before handling: 42786
Missing values handled using forward fill.
Total missing values after handling: 9

```

C:\Users\rahul\AppData\Local\Temp\ipykernel_8304\1573965327.py:3:
FutureWarning: DataFrame.fillna with 'method' is deprecated and will
raise in a future version. Use obj.ffill() or obj.bfill() instead.
    full_df.fillna(method='ffill', inplace=True)

# =====
# Step 3: Initial Analysis and Reporting
# =====
print("\n--- Ophthalmology Drug Tracking Analysis ---")

# 1. Identified and report on unscanned vials
unscanned_vials = full_df[full_df['Unscanned'] == True]
```

```

print(f"\n## Unscanned Vials Report ({len(unscanned_vials)} found)")
if not unscanned_vials.empty:
    print("The following vials were dispensed but not scanned, leading
to potential revenue loss:")
    print(unscanned_vials[['Vial Number', 'Lot Number', 'Purchase
Price', 'Dispense ID', 'Patient ID', 'Date Of Dispense',
'Username']].head())

```

--- Ophthalmology Drug Tracking Analysis ---

```

## Unscanned Vials Report (875 found)
The following vials were dispensed but not scanned, leading to
potential revenue loss:
      Vial Number  Lot Number  Purchase Price  Dispense ID  Patient
ID \
6   VN-0000005   LOT-8583121          50.0  DISP-010328  567386.0
34  VN-0000025   LOT-5251569          50.0  DISP-003474  930415.0
35  VN-0000026   LOT-8694199          50.0  DISP-003474  930415.0
50  VN-0000038   LOT-1985938          50.0  DISP-002492  435402.0
71  VN-0000050   LOT-8239582          50.0  DISP-007851  722927.0

```

	Date Of Dispense	Username
6	2025-05-04	jen.kirby
34	2025-03-01	dina.dropem
35	2025-03-01	dina.dropem
50	2025-05-08	dina.dropem
71	2025-06-11	sarah.benz

```

# 2. Identified and report on denied claims
denied_claims = full_df[full_df['Denied'] == True]
print(f"\n## Denied Claims Report ({len(denied_claims)} found)")
if not denied_claims.empty:
    print("The following claims were denied by insurance:")
    print(denied_claims[['Claim ID', 'Dispense ID', 'Vial Number',
'Patient ID', 'Claim Date', 'Amount']].head())

```

```

## Denied Claims Report (561 found)
The following claims were denied by insurance:
      Claim ID  Dispense ID  Vial Number  Patient ID  Claim Date
Amount
6     CLM-010328  DISP-010328  VN-0000005   567386.0  2025-09-09
92.40
50    CLM-002492  DISP-002492  VN-0000038   435402.0  2025-03-26
102.98

```

```

71 CLM-007851 DISP-007851 VN-0000050 722927.0 2025-02-22
119.08
77 CLM-000043 DISP-000043 VN-0000056 702136.0 2025-05-07
65.80
201 CLM-008965 DISP-008965 VN-0000149 406197.0 2025-07-21
124.17

# 3. Performed and report on profitability analysis
reimbursed_claims = full_df[(full_df['Denied'] == False) &
(full_df['Amount'].notna())].copy()
reimbursed_claims['Profit'] = reimbursed_claims['Amount'] -
reimbursed_claims['Purchase Price']
total_revenue = reimbursed_claims['Amount'].sum()
total_cost = reimbursed_claims['Purchase Price'].sum()
total_profit = reimbursed_claims['Profit'].sum()

print("\n## Profitability Analysis")
print(f"Total Revenue from Reimbursed Claims: ${total_revenue:,.2f}")
print(f"Total Cost of Goods Sold: ${total_cost:,.2f}")
print(f"Total Profit: ${total_profit:,.2f}")

## Profitability Analysis
Total Revenue from Reimbursed Claims: $1,980,313.37
Total Cost of Goods Sold: $759,600.00
Total Profit: $1,220,713.37

#
=====
=====

# Step 4: Exploratory Data Analysis (EDA) and Visualizations
#
=====

print("\n--- Generating EDA Visualizations ---")

# Visualization 1: Distribution of Key Numerical Features
plt.figure(figsize=(14, 6))
plt.suptitle('Distribution of Numerical Features')

plt.subplot(1, 2, 1)
sns.histplot(full_df['Purchase Price'], kde=True, bins=30)
plt.title('Vial Purchase Price')
plt.xlabel('Purchase Price ($)')
plt.ylabel('Frequency')

plt.subplot(1, 2, 2)
sns.histplot(full_df[full_df['Amount'].notna()]['Amount'], kde=True,
bins=30, color='green')
plt.title('Claim Reimbursement Amount')

```


```

plt.xlabel('Reimbursement Amount ($)')
plt.ylabel('Frequency')

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

```


--- Generating EDA Visualizations ---




```

# Visualization 2: Relationship between Purchase Price and
# Reimbursement Amount
plt.figure(figsize=(8, 6))
sns.scatterplot(data=full_df, x='Purchase Price', y='Amount',
hue='Denied')
plt.title('Purchase Price vs. Reimbursement Amount')
plt.xlabel('Purchase Price ($)')
plt.ylabel('Reimbursement Amount ($)')
plt.legend(title='Claim Denied')
plt.show()

```



```
# Visualization 3: Correlation Heatmap
plt.figure(figsize=(10, 8))
numeric_cols = full_df.select_dtypes(include=['float64',
'int64']).columns
correlation_matrix = full_df[numeric_cols].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f")
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```



```

#
=====
=====

# Step 5: Hypothesis Testing and Feature Engineering
#
=====

# Hypothesis Test: Are reimbursement amounts different for approved
vs. denied claims?
approved_claims_amount = full_df[(full_df['Denied'] == False) &
(full_df['Amount'].notna())]['Amount']
denied_claims_amount = full_df[(full_df['Denied'] == True) &
(full_df['Amount'].notna())]['Amount']

t_stat, p_value = stats.ttest_ind(approved_claims_amount,

```

```

denied_claims_amount, equal_var=False, nan_policy='omit')

print("\n--- Hypothesis Test: Approved vs. Denied Claim Amounts ---")
print(f"T-statistic: {t_stat:.4f}")
print(f"P-value: {p_value:.4f}")

--- Hypothesis Test: Approved vs. Denied Claim Amounts ---
T-statistic: 0.0213
P-value: 0.9831

alpha = 0.05
if p_value < alpha:
    print("\nConclusion: We reject the null hypothesis.")
    print("There is a statistically significant difference in
reimbursement amounts.")
else:
    print("\nConclusion: We fail to reject the null hypothesis.")
    print("There is no statistically significant difference in
reimbursement amounts.")

Conclusion: We fail to reject the null hypothesis.
There is no statistically significant difference in reimbursement
amounts.

# Feature Engineering: Adding 'Profit Margin'
reimbursed_claims['Profit Margin'] = (reimbursed_claims['Profit'] /
reimbursed_claims['Amount']) * 100


print("\n--- Feature Engineering: Profit Margin ---")
print("Top 5 Vials by Profit Margin:")
print(reimbursed_claims.nlargest(5, 'Profit Margin')[['Vial Number',
'Purchase Price', 'Amount', 'Profit', 'Profit Margin']])


--- Feature Engineering: Profit Margin ---
Top 5 Vials by Profit Margin:
   Vial Number Purchase Price   Amount   Profit  Profit Margin
15      VN-0000010       50.0  199.98  149.98    74.997500
16      VN-0000011       50.0  199.98  149.98    74.997500
17      VN-0000012       50.0  199.98  149.98    74.997500
2695     VN-0002064       50.0  199.96  149.96    74.994999
13870    VN-0010556       50.0  199.94  149.94    74.992498

plt.figure(figsize=(8, 6))
sns.histplot(reimbursed_claims['Profit Margin'], kde=True, bins=30,
color='purple')
plt.title('Distribution of Profit Margin')
plt.xlabel('Profit Margin (%)')

```

```
plt.ylabel('Frequency')
plt.show()
```



```
#
=====
=====
# Step 6: Creating a Train-Test Split for Modeling
#
=====
```

print("\n--- Preparing Data for Modeling ---")

```
# 1. Defining the target variable (y)
# It will predict if a claim is denied (1) or not (0).
# Filling NaN with 0
full_df['Denied'] = full_df['Denied'].fillna(0).astype(int)
y = full_df['Denied'].astype(int)
```

```

--- Preparing Data for Modeling ---

# 2. Selecting features (X) to be used for prediction.
# Dropped identifiers and columns that are not useful for prediction.
features_to_use = full_df.drop(columns=[
    'Vial Number', 'Lot Number', 'Expiration Date', 'Invoice Number',
    'Invoice Date', 'Dispense ID', 'Patient ID', 'Date Of Dispense',
    'Claim ID', 'Claim Date', 'Denied', 'Unscanned'
])

# 3. One-hot encode the categorical features to convert them to numbers.
X = pd.get_dummies(features_to_use, drop_first=True)

# 4. Performed the Train-Test Split (80% train, 20% test)
# 'stratify=y' ensuring the proportion of denied claims is the same in both sets.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 5. Verified the split and show the shapes of the new datasets
print("\nData successfully split for modeling.")
print("--- Shape of the Datasets ---")
print(f"Training Features (X_train): {X_train.shape}")
print(f"Testing Features (X_test): {X_test.shape}")
print(f"Training Target (y_train): {y_train.shape}")
print(f"Testing Target (y_test): {y_test.shape}")

print("\n--- Distribution of Target in Train vs. Test Sets ---")
print("Training Set (Denied=1):\n",
y_train.value_counts(normalize=True))
print("\nTesting Set (Denied=1):\n",
y_test.value_counts(normalize=True))

print("\n\n--- Conclusion ---")
print("The datasets (X_train, X_test, y_train, y_test) are now ready for model building.")

```

Data successfully split for modeling.
--- Shape of the Datasets ---
Training Features (X_train): (12603, 9)
Testing Features (X_test): (3151, 9)
Training Target (y_train): (12603,)
Testing Target (y_test): (3151,)

--- Distribution of Target in Train vs. Test Sets ---
Training Set (Denied=1):
Denied

```
0      0.964374
1      0.035626
Name: proportion, dtype: float64

Testing Set (Denied=1):
 Denied
0      0.964456
1      0.035544
Name: proportion, dtype: float64

--- Conclusion ---
The datasets (X_train, X_test, y_train, y_test) are now ready for
model building.
```